

Tesina

Iñaki Garay

Septiembre 2020

Contents

Introducción / Esquema General

- Los compiladores tradicionales tienen una arquitectura de pipeline.
- Los editores y entornos de desarrollo modernos usan LSP (Language Server Protocol). Porque?
- Una implementación de LSP requiere de componentes de compiladores (especialmente análisis).
- La arquitectura de pipeline no se adapta bien a estos requerimientos modernos de reducir tiempos de compilación y de proveer información online durante edición.
- Estos objetivos pueden ser logrados mediante compilación incremental.
- La compilación incremental puede ser lograda mediante una arquitectura basada en queries.
- La arquitectura basada en queries es implementada tomando inspiración de build systems.
- Rust-analyzer es la segunda implementación de LSP para Rust.
- Porque no funciona la primera iteración?
- Rust-Analyzer usa una librería, Salsa, para cachear las queries parciales.
- Como funciona Salsa?

Motivaciones

Los compiladores ya no son cajas negras que ingieren un conjunto de archivos fuente y producen código ensamblador. De compiladores modernos se espera que:

- Sean incrementales, es decir, si se recompila el proyecto después de producir modificaciones en el código fuente, solo se recompila lo que fue afectado por esas modificaciones.

- Provean funcionalidad para editores, e.g. saltar a definición, encontrar el tipo de una expresión en una ubicación dada, y mostrar errores al editar.

Una arquitectura de compilador tradicional se
GRAFICO

- We push source text down a pipeline and run a fixed set of transformations until we finally output assembly code or some other target language. Along the way we often need to read and update some state. For example, we might update a type table during type checking so we can later look up the type of entities that the code refers to.

Language Server Protocol

The Language Server Protocol (LSP) is an open, JSON-RPC-based protocol for use between source code editors or integrated development environments (IDEs) and servers that provide programming language-specific features. The goal of the protocol is to allow programming language support to be implemented and distributed independently of any given editor or IDE.

Adding features like auto complete, go to definition, or documentation on hover for a programming language takes significant effort. Traditionally this work had to be repeated for each development tool, as each tool provides different APIs for implementing the same feature.

A Language Server is meant to provide the language-specific smarts and communicate with development tools over a protocol that enables inter-process communication.

The idea behind the Language Server Protocol (LSP) is to standardize the protocol for how such servers and development tools communicate. This way, a single Language Server can be re-used in multiple development tools, which in turn can support multiple languages with minimal effort.

LSP is a win for both language providers and tooling vendors!

Modern integrated development environments (IDEs) provide developers with sophisticated features like code completion, refactoring, navigating to a symbol's definition, syntax highlighting, and error and warning markers.

For example, in a text-based programming language, a programmer might want to rename a method `read`. The programmer could either manually edit the respective source code files and change the appropriate occurrences of the old method name into the new name, or instead use an IDE's refactoring capabilities to make all the necessary changes automatically. To be able to support this style of refactoring, an IDE needs a sophisticated understanding of the programming language that the program's source is written in. A programming tool without such an understanding—for example, one that performs a naive search-and-replace instead—could introduce errors. When renaming a `read` method, for example, the tool should not replace the partial match in a variable that might be called `readyState`, nor should it replace the portion of a code comment containing the word "already". Neither should renaming a local variable `read`, for example, end up altering similarly named variables in other scopes.

Conventional compilers or interpreters for a specific programming language are typically unable to provide these language services, because they are written with the goal of either transforming the source code into object code or immediately executing the code. Additionally, language services must be able to handle source code that is not well-formed, e.g. because the programmer is in the middle of editing and has not yet finished typing a statement, procedure, or other construct. Additionally, small changes to a source code file which are done during typing usually change the semantics of the program. In order to provide instant feedback to the user, the editing tool must be able to very quickly evaluate the syntactical and semantical consequences of a specific modification. Compilers and interpreters therefore provide a poor candidate for producing the information needed for an editing tool to consume.

Prior to the design and implementation of the Language Server Protocol for the development of Visual Studio Code, most language services were generally tied to a given IDE or other editor. In the absence of the Language Server Protocol, language services are typically implemented by utilizing a tool-specific extension API. Providing the same

language service to another editing tool requires effort to adapt the existing code so that the service may target the second editor's extension interfaces.

The Language Server Protocol allows for decoupling language services from the editor so that the services may be contained within a general purpose language server. Any editor can inherit sophisticated support for many different languages by making use of existing language servers. Similarly, a programmer involved with the development of a new programming language can make services for that language available to existing editing tools. Making use of language servers via the Language Server Protocol thus also reduces the burden on vendors of editing tools, because vendors do not need to develop language services of their own for the languages the vendor intends to support, as long as the language servers have already been implemented. The Language Server Protocol also enables the distribution and development of servers contributed by an interested third-party, such as end users, without additional involvement by either the vendor of the compiler for the programming language in use or the vendor of the editor to which the language support is being added.

LSP is not restricted to programming languages. It can be used for any kind of text-based language, like specifications or domain-specific languages (DSL).

Velocidad de Compilación

- Improving compile times has actually been a major development focus after Rust reached 1.0 -- although, up to this point, much of the work towards this goal has gone into laying architectural foundations within the compiler and we are only slowly beginning to see actual results.
- One of the projects that is building on these foundations, and that should help improve compile times a lot for typical workflows, is incremental compilation.
Incremental compilation avoids redoing work when you recompile a crate, which will ultimately lead to a much faster edit-compile-debug cycle.

Mecanismos

Compilación Incremental

La compilación incremental es una forma de computación incremental aplicada a la compilación. En contraste con compiladores comunes que realizan "builds limpios" y ante un cambio en el código fuente recompilan todas las unidades de compilación, un compilador incremental solo recompila las unidades modificadas. Al construir sobre el trabajo hecho previamente, el compilador incremental evita la ineficiencia de repetir trabajo ya realizado. Se puede decir que un compilador incremental reduce la granularidad de las unidades de compilación tradicionales a la vez que mantiene la semántica del lenguaje.

Muchas herramientas de desarrollo aprovechan compiladores incrementales para proveer a sus usuarios un entorno mucho mas interactivo. No es inusual que un compilador incremental sea invocado por cada cambio en un archivo fuente, de tal manera que el usuario es informado inmediatamente de cualquier error de compilación causado por sus modificaciones. Este esquema, en contraste con el modelo de compilación tradicional, acorta el ciclo de desarrollo considerablemente.

Una desventaja de este esquema es que el compilador no puede optimizar fácilmente el código que compila, dada la localidad y el alcance reducido de los cambios. Normalmente esto no es un problema, dado que la optimización del código generado se aplica solamente al producir un *release build*, instancia en la cual se puede usar el compilador tradicional.

Arquitectura Basada en Consultas

- Going from pipeline to queries
- What does it take to get the type of a qualified name?
- In a pipeline-based architecture we would just look it up in the type table.
- With queries, we have to think differently.
- Instead of relying on having updated some piece of state, we do it as if it was done from scratch.
- As a first iteration, we do it completely from scratch

- We first find out what file the name comes from, which might be `Data/List.vix` for `Data.List`, then read the contents of the file, parse it, perhaps we do name resolution to find out what the names in the code refer to given what is imported, and last we look up the name-resolved definition and type check it, returning its type.

```
-
'''
fetchType :: QualifiedName -> IO Type
fetchType (QualifiedName moduleName name) = do
  fileName <- moduleFileName moduleName
  sourceCode <- readFile fileName
  parsedModule <- parseModule sourceCode
  resolvedModule <- resolveNames parsedModule
  let definition = lookup name resolvedModule
  inferDefinitionType definition
'''
```

- Let's first refactor the code into smaller functions:
- Note that each of the functions do everything from scratch on their own, i.e. they're each doing a (longer and longer) prefix of the work you'd do in a pipeline.
- I've found this to be a common pattern in my query-based compilers.
- One way to make this efficient would be to add a memoisation layer around each function.
- That way, we do some expensive work the first time we invoke a function with a specific argument, but subsequent calls are cheap as they can return the cached result.
- This is essentially what we'll do, but we won't use a separate cache per function, but instead have a central cache, indexed by the query.
- **Why Incremental Compilation in the First Place?**
- Much of a programmer's time is spent in an edit-compile-debug workflow:
- you make a small change (often in a single module or even function),

- you let the compiler translate the code into a binary, and finally
- you run the program or a bunch of unit tests in order to see results of the change.
- After that it's back to step one, making the next small change informed by the knowledge gained in the previous iteration.
This essential feedback loop is at the core of our daily work.
We want the time being stalled while waiting for the compiler to produce an executable program to be as short as possible.
- Incremental compilation is a way of exploiting the fact that little changes between compiles during the regular programming workflow:
Many, if not most, of the changes done in between two compilation sessions only have local impact on the machine code in the output binary, while the rest of the program, same as at the source level, will end up exactly the same, bit for bit.
Incremental compilation aims at retaining as much of those unchanged parts as possible while redoing only that amount of work that actually must be redone.
- ****How Do You Make Something "Incremental"?****
- We have already heard that computing something incrementally means updating only those parts of the computation's output that need to be adapted in response to a given change in the computation's inputs.
- One basic strategy we can employ to achieve this is to view one big computation (like compiling a program) as a composite of many smaller, interrelated computations that build up on each other.
- Each of those smaller computations will yield an intermediate result that can be cached and hopefully re-used in a later iteration, sparing us the need to re-compute that particular intermediate result again.
- ****An Incremental Compiler****
- The way we chose to implement incrementality in the Rust compiler is straightforward: An incremental compilation session follows exactly the same steps in the same order as a batch compilation session.
- However, when control flow reaches a point where it is about to compute some non-trivial intermediate result, it will try to load that result from the incremental compilation cache on disk instead.
- If there is a valid entry in the cache, the compiler can just skip computing that particular piece of data. Let's take a look at a (simplified) overview of the different compilation phases and the

intermediate results they produce:

- First the compiler will parse the source code into an abstract syntax tree (AST).

The AST then goes through the analysis phase which produces type information and the MIR for each function.

After that, if analysis did not find any errors, the codegen phase will transform the MIR version of the program into its machine code version, producing one object file per source-level module.

In the last step all the object files get linked together into the final output binary which may be a library or an executable.

- So, this seems pretty simple so far: Instead of computing something a second time, just load the value from the cache.

Things get tricky though when we need to find out if it's actually valid to use a value from the cache or if we have to re-compute it because of some changed input.

Memoization y Seguimiento de Dependencias

- Rock, Shake, Salsa
- This functionality is provided by Rock, a library that packages up some functionality for creating query-based compilers.
- Rock is an experimental library heavily inspired by Shake and the Build systems à la carte paper.
- It essentially implements a build system framework, like make.
- Build systems have a lot in common with modern compilers since we want them to be incremental, i.e. to take advantage of previous build results when building anew with few changes.
- But there's also a difference: Most build systems don't care about the types of their queries since they work at the level of files and file systems.
- Build systems à la carte is closer to what we want.
- There the user writes a bunch of computations, tasks, choosing a suitable type for keys and a type for values.
- The tasks are formulated assuming they're run in an environment where there is a function `fetch` of type `Key -> Task Value`, where `Task` is a type for describing build system rules, that can be used to fetch the value of a dependency with a specific key.
- In our above example, the key type might look like this:
- The build system has control over what code runs when we do a `fetch`, so by varying that it can do fine-grained dependency tracking, memoisation, and incremental updates.

- Build systems à la carte is also about exploring what kind of build systems we get when we vary what Task is allowed to do, e.g. if it's a Monad or Applicative.
- In Rock, we're not exploring that, so our Task is a thin layer on top of IO.
- A problem that pops up now, however, is that there's no satisfactory type for Value.
- We want `fetch (ParsedModuleKey "Data.List")` to return a `ParsedModule`, while `fetch (TypeKey "Data.List.map")` should return something of type `Type`.
- Indexed queries

- Rock allows us to index the key type by the return type of the query. The Key type in our running example becomes the following GADT:

```
'''
```

```
data Key a where
```

```
  ParsedModuleKey :: ModuleName -> Key ParsedModule
```

```
  ResolvedModuleKey :: ModuleName -> Key ResolvedModule
```

```
  TypeKey :: QualifiedName -> Key Type
```

```
'''
```

- The `fetch` function gets the type forall `a. Key a -> Task a`, so we get a `ParsedModule` when we run `fetch (ParsedModuleKey "Data.List")`, like we wanted, because the return type depends on the key we use.
- Now that we know what `fetch` should look like, it's also worth revealing what the Task type looks like in Rock, more concretely.
- As mentioned, it's a thin layer around IO, providing a way to fetch keys (like Key above):
- The rules of our compiler, i.e. its "Makefile", then becomes the following function, reusing the functions from above:

```
'''
```

```
rules :: Key a -> Task a
```

```
rules key = case key of
```

```
  ParsedModuleKey moduleName ->
```

```
    fetchParsedModule moduleName
```

```
  ResolvedModuleKey moduleName ->
```

```
    fetchResolvedModule moduleName
```

```
  TypeKey qualifiedName ->
```

```
    fetchType qualifiedName
```

```
'''
```

- Caching

- The most basic way to run a Task in Rock is to directly call the rules function when a Task fetches a key.
- This results in an inefficient build system that recomputes every query from scratch.
- But the Rock library lets us layer more functionality onto our rules function, and one thing that we can add is memoisation.
- If we do that Rock caches the result of each fetched key by storing the key-value pairs of already performed fetches in a dependent hashmap.
- This way, we perform each query at most once during a single run of the compiler.
- Verifying dependencies and reusing state
- Another kind of functionality that can be layered onto the rules function is incremental updates. When it's used, Rock keeps track of what dependencies a task used when it was executed (much like Shake) in a table, i.e. what keys it fetched and what the values were.
- Using this information it's able to determine when it's safe to reuse the cache from a previous run of the compiler even though there might be changes in other parts of the dependency graph.
- This fine-grained dependency tracking also allows reusing the cache when a dependency of a task changes in a way that has no effect.
- For example, whitespace changes might trigger a re-parse, but since the AST is the same, the cache can be reused in queries that depend on the parse result.
- Reverse dependency tracking
- Verifying dependencies can be too slow for real-time tooling like language servers, because large parts of the dependency graph have to be traversed just to check that most of it is unchanged even for tiny changes.
- For example, if we make changes to a source file with many large imports, we need to walk the dependency trees of all of the imports just to update the editor state for that single file.
- This is because dependency verification by itself needs to go all the way to the root queries for all the dependencies of a given query, which can often be a large proportion of the whole dependency tree.
- To fix this, Rock can also be made to track reverse dependencies

- between queries.
- When e.g. a language server detects that a single file has changed, the reverse dependency tree is used to invalidate the cache just for the queries that depend on that file by walking the reverse dependencies starting from the changed file.
 - Since the imported modules don't depend on that file, they don't need to be re-checked, resulting in much snappier tooling!
- ****Dependency Graphs****
- There is a formal method that can be used to model a computation's intermediate results and their individual "up-to-dateness" in a straightforward way: dependency graphs.
 - It looks like this: Each input and each intermediate result is represented as a node in a directed graph.
The edges in the graph, on the other hand, represent which intermediate result or input can have an influence on some other intermediate result.
 - Note, by the way, that the above graph is a tree just because the computation it models has the form of a tree.
In general, dependency graphs are directed acyclic graphs
 - What makes this data structure really useful is that we can ask it questions of the form "if X has changed, is Y still up-to-date?".
We just take the node representing Y and collect all the inputs Y depends on by transitively following all edges originating from Y.
If any of those inputs has changed, the value we have cached for Y cannot be relied on anymore.
- ****Dependency Tracking in the Compiler****
- When compiling in incremental mode, we always build the dependency graph of the produced data: every time, some piece of data is written (like an object file), we record which other pieces of data we are accessing while doing so.
 - The emphasis is on recording here. At any point in time the compiler keeps track of which piece of data it is currently working on (it does so in the background in thread-local memory).
 - This is the currently active node of the dependency graph.
Conversely, the data that needs to be read to compute the value of the active node is also tracked: it usually already resides in some kind container (e.g. a hash table) that requires invoking a lookup method

- to access a specific entry.
- We make good use of this fact by making these lookup methods transparently create edges in the dependency graph: whenever an entry is accessed, we know that it is being read and we know what it is being read for (the currently active node).
 - This gives us both ends of the dependency edge and we can simply add it to the graph.
At the end of the compilation sessions we have all our data nicely linked up, mostly automatically.
 - This dependency graph is then stored in the incremental compilation cache directory along with the cache entries it describes.
 - At the beginning of a subsequent compilation session, we detect which inputs (=AST nodes) have changed by comparing them to the previous version.
Given the graph and the set of changed inputs, we can easily find all cache entries that are not up-to-date anymore and just remove them from the cache.
 - Anything that has survived this cache validation phase can safely be re-used during the current compilation session.
 - There are a few benefits to the automated dependency tracking approach we are employing.
Since it is built into the compiler's internal APIs, it will stay up-to-date with changes to the compiler, and it is hard to accidentally forget about.
And if one still forgets using it correctly (e.g. by not declaring the correct active node in some place) then the result is an overly conservative, but still "correct" dependency graph:
It will negatively impact the re-use ratio but it will not lead to incorrectly re-using some outdated piece of data.
 - Another aspect is that the system does not try to predict or compute what the dependency graph is going to look like, it just keeps track.
A large part of our (yet to be written) regression tests, on the other hand, will give a description of what the dependency graph for a given program ought to look like. This makes sure that the actual graph and the reference graph are arrived at via different methods, reducing the risk that both the compiler and the test case agree on the same, yet wrong, value.
 - ****"Faster! Up to 15% or More."****
 - Let's take a look at some of the implications of what we've learned so

far:

- The dependency graph reflects the actual dependencies between parts of the source code and parts of the output binary.
- If there is some input node that is reachable from many intermediate results, e.g. a central data type that is used in almost every function, then changing the definition of that data type will mean that everything has to be compiled from scratch, there's no way around it.
- In other words, the effectiveness of incremental compilation is very sensitive to the structure of the program being compiled and the change being made. Changing a single character in the source code might very well invalidate the whole incremental compilation cache. Usually though, this kind of change is a rare case and most of the time only a small portion of the program has to be recompiled.

Caso de Estudio: Rustc

Rustc, el compilador de rust, tiene su propia implementacion de queries.

- ****The Current Status of the Implementation**** (09/2019)
- For the first spike implementation of incremental compilation, what we call the alpha version now, we chose to focus on caching object files.
- Consequently, if this phase can be skipped at least for part of a code base, this is where the biggest impact on compile times can be achieved.
- With that in mind, we can also give an upper bound on how much time this initial version of incremental compilation can save:
If the compiler spends X seconds optimizing when compiling your crate, then incremental compilation will reduce compile times at most by those X seconds.
- Another area that has a large influence on the actual effectiveness of the alpha version is dependency tracking granularity: It's up to us how fine-grained we make our dependency graphs, and the current implementation makes it rather coarse in places.
For example, the dependency graph only knows a single node for all methods in an impl. As a consequence, the compiler will consider all methods of that impl as changed if just one of them is changed. This of course will mean that more code will be re-compiled than is strictly necessary.

Caso de Estudio: Rust-analyzer y Salsa

Rust-Analyzer utiliza una librería llamada salsa.

Caso de Estudio: Como funciona salsa?

La idea central de salsa es definir el programa como un conjunto de *queries*. Cada query se usa como una función $K \rightarrow V$ que mapea de una clave de tipo K a un valor de tipo V .

Las queries en salsa son de dos variedades básicas:

- **Entradas:** definen los inputs básicos al sistema, los cuales pueden cambiar en cualquier momento.
- **Funciones:** funciones puras (sin efectos secundarios) que transforman las entradas en otros valores. Los resultados de estas queries se memoizan para evitar recomputarlas. Cuando se modifican las entradas, salsa determina cuáles valores memoizados pueden ser reutilizados y cuáles deben ser recomputados.

El esquema general de utilización de salsa consiste en tres pasos:

1. Definir uno o más grupos de queries que contendrán las entradas y las queries requeridas. Se puede definir más de un grupo para separar las queries en componentes.
2. Definir las queries.
3. Definir la base de datos, la cual contendrá el almacenamiento para las entradas y queries utilizadas.

Conclusiones

- Most modern languages need to have a strategy for tooling, and building compilers around query systems seems like an extremely promising approach to me.
- With queries the compiler writer doesn't have to handle updates to and invalidation of a bunch of ad-hoc caches, which can be the result when adding incremental updates to a traditional compiler pipeline.
- In a query-based system it's all handled centrally once and for all, which means there's less of a chance it's wrong.
- Queries are excellent for tooling because they allow us to ask for the value of any query at any time without worrying about order or temporal effects, just like a well-written Makefile.
- The system will compute or retrieve cached values for the query and

- its dependencies automatically in an incremental way.
- Query-based compilers are also surprisingly easy to parallelise.
 - Since we're allowed to make any query at any time, and they're memoised the first time they're run, we can fire off queries in parallel without having to think much.
 - In Sixty, the default behaviour is for all input modules to be type checked in parallel.
 - **Future Plans** (09/2019)
 - The section on the current status already laid out the two major axes along which we will pursue increased efficiency:
 - Cache more intermediate results, like MIR and type information, which will allow the compiler to skip more and more steps.
 - Make dependency tracking more precise, so that the compiler encounters fewer false positives during cache invalidation.

Fuentes y Referencias

- General
 - [✓ Youtube: Anders Hejlsberg on Modern Compiler Construction](#)
 - [✓ Wikipedia: Incremental Compiler](#)
 - [✓ Olle Fredriksson: Query-based compiler architectures](#)
 - [✓ Rust Blog: Incremental Compilation](#)
 - [□ Build Systems A La Carte](#)
 - [□ Youtube: 2016 LLVM Developers' Meeting: D. Dunbar "A New Architecture for Building Software"](#)
 - [□ An approach to incremental compilation](#)
- Rustc Dev Guide
 - [□ Overview of the Compiler](#)
 - [□ High-level overview of the compiler source](#)
 - [□ Queries: demand-driven compilation](#)
 - * [□ The Query Evaluation Model in Detail](#)
 - * [□ Incremental compilation](#)
 - * [□ Incremental Compilation In Detail](#)
 - * [□ Debugging and Testing Dependencies](#)
 - * [□ Profiling Queries](#)
 - * [□ How Salsa works](#)
- Rust Analyzer
 - [□ Rust Analyzer](#)

- [☐Manual](#)
- [☐Blog](#)
- [☐rust-analyzer/tree/master/docs/dev](#)
- [☐Rust Analyzer in 2018 and 2019](#)
- [☐Status of rust-analyzer](#)
- [☐2020 Intro to Rust Analyzer](#)
- [☐2020 What I learned contributing to Rust-Analyzer](#)
- [☐Youtube: Are we *actually* IDE yet? A look on the Rust IDE Story - Igor Matuszewski](#)
- [☐Youtube: Rust analyzer guide](#)
- [☐Youtube: rust analyzer syntax trees](#)
- [☐Youtube: rust-analyzer type-checker overview by flodiebold](#)
- [☐Youtube: Rust Analyzer Q&A](#)
- Salsa
 - [☐The Salsa Book](#)
 - [☐Youtube: Incremental Compilation Working Group](#)
 - [☐Youtube: Responsive compilers - Nicholas Matsakis - PLISS 2019](#)
 - [☐Youtube: Things I Learned \(TIL\) - Nicholas Matsakis - PLISS 2019](#)
 - [☐Youtube: How Salsa Works \(2019.01\)](#)
 - [☐Youtube: Salsa In More Depth \(2019.01\)](#)
 - [☐Youtube: RLS 2.0, Salsa, and Name Resolution](#)
- Rust Compilation Speed
 - [☐How to alleviate the pain of Rust compile times](#)
 - [☐Nethercote: How to speed up the Rust compiler](#)
 - [☐Nethercote: How to speed up the Rust compiler some more](#)
 - [☐Nethercote: How to speed up the Rust compiler in 2018](#)
 - [☐Nethercote: How to speed up the Rust compiler some more in 2018](#)
 - [☐Nethercote: How to speed up the Rust compiler in 2018: NLL edition](#)
 - [☐Nethercote: The Rust compiler is getting faster](#)
 - [☐Nethercote: The Rust compiler is still getting faster](#)
 - [☐Nethercote: How to speed up the Rust compiler in 2019](#)
 - [☐Nethercote: How to speed up the Rust compiler some more in 2019](#)
 - [☐Nethercote: How to speed up the Rust compiler one last time in 2019](#)

- [□Nethercote: How to speed up the Rust compiler in 2020](#)
- [□Nethercote: How to speed up the Rust compiler some more in 2020](#)
- [□Nethercote: How to speed up the Rust compiler one last time](#)
- [□PingCAP Blog: The Rust Compilation Model Calamity](#)
- [□PingCAP Blog: Generics and Compile-Time in Rust](#)
- [□PingCAP Blog: Rust’s Huge Compilation Units](#)
- [□PingCAP Blog: A Few More Reasons Rust Compiles Slowly](#)
- Miscellaneous
 - [□Youtube: Making Fast Incremental Compiler for Huge Codebase - Michał Bartkowiak - code::dive 2019](#)
 - [□Youtube: Starting with Semantics - Sylvan Clebsch - PLISS 2019](#)
 - [□Youtube: Polyhedral Compilation as a Design Pattern for Compilers \(1/2\) - Albert Cohen - PLISS 2019](#)
 - [□Youtube: Polyhedral Compilation as a Design Pattern for Compilers \(2/2\) - Albert Cohen - PLISS 2019](#)
 - [□Youtube: First-Class Continuations: What and Why - Arjun Guha](#)
 - [□Youtube: Implementing First-Class Continuations by Source to Source Translation - Arjun Guha - PLISS 2019](#)
 - [□Youtube: Static Program Analysis \(part 1/2\) - Anders Møller - PLISS 2019](#)
 - [□Youtube: Static Program Analysis \(part 2/2\) - Anders Møller - PLISS 2019](#)